

Statistics

Dr. Polla Fattah

What You Will Learn

By the end of this crash course, you'll be able to:

- Summarize and visualize data
- Measure variability and shape
- Construct and interpret confidence intervals
- Conduct hypothesis testing (t-tests, ANOVA, chi-square)
- Analyze categorical and numeric outcomes
- Choose appropriate statistical tests for your thesis
- Communicate statistical results effectively

Key Questions Every Researcher Must Answer

- Does my algorithm really perform better than others?
- How confident can I be in my results?
- Is the difference I observed due to chance?
- How do I report my findings objectively?

Statistics provides the **language of evidence** in science.

Why Postgraduates Can't Ignore Statistics

Without statistical analysis, a statement like:

"I ran my new sorting algorithm and it seemed faster than QuickSort."
is anecdotal at best.

In a thesis defense, your examiners will ask:

- How many times did you test it?
- What was the average run-time?
- How much variability was there?
- Could this be due to randomness?

- Was your test design valid?

These are not programming questions — they are **statistical**.

Statistics in Experimental Design and Validation

Activity	Statistical Role
Comparing algorithms	t-tests, ANOVA
Classifier performance	Confidence intervals
Survey data	Chi-square, proportions
Benchmarking systems	Regression
System reliability	Error rate, standard deviation

Good experimental design includes:

- Stating clear hypotheses
- Designing fair trials
- Collecting and summarizing results
- Performing appropriate statistical analysis

Our Guiding Scenario: The Algorithm Showdown

To make this course practical, we'll follow one running example:

You're defending your MSc thesis:

"Comparing MergeSort and QuickSort performance."

You've:

- Run each algorithm 30 times
- Recorded run-time (ms)
- Collected performance data

Your core research question:

> *"Is MergeSort significantly faster than QuickSort?"*

We will use this example throughout to introduce and apply statistical concepts.

Descriptive Statistics: Center and Spread

To begin answering our thesis question (“Is MergeSort faster than QuickSort?”), we need to **describe** the data.

We'll start with:

- Measures of central tendency
- Measures of variability
- How they behave in symmetric and skewed data
- Practical examples with our simulated results

Understanding Central Tendency

Central tendency describes the “middle” of the data — where values tend to cluster.

Three common measures:

- **Mean:** Arithmetic average
- **Median:** Middle value when sorted
- **Mode:** Most frequent value (less useful in continuous data)

When Do These Measures Differ?

In a perfectly **symmetric distribution**, all three measures are equal:

$\text{Mean} = \text{Median} = \text{Mode}$

In a **right-skewed** distribution (tail on the right):

$\text{Mean} > \text{Median} > \text{Mode}$

In a **left-skewed** distribution (tail on the left):

$\text{Mean} < \text{Median} < \text{Mode}$

This helps analysts infer the **shape** of the data from summary statistics.

Simulating Our Data (QuickSort vs MergeSort)

We'll work with the following data throughout this presentation:

```
set.seed(42)
qs <- rnorm(30, mean = 52, sd = 6)
ms <- rnorm(30, mean = 48, sd = 5)

algo_data <- data.frame(
  Algorithm = rep(c("QuickSort", "MergeSort"), each =
    30),
  Time_ms = c(qs, ms)
)
```

Each row represents one run of either algorithm. We'll use this data in visualizations and statistical analysis.

Displaying the Data as a Table

It's useful to see the actual values used:

Trial	QuickSort	MergeSort
1	60.23	50.28
2	48.61	51.52
3	54.18	53.18
4	55.80	44.96
5	54.43	50.52
6	51.36	39.41
7	61.07	44.08
8	51.43	43.75
9	64.11	35.93

Trial	QuickSort	MergeSort
10	51.62	48.18

Mean and Median of Each Algorithm

```
mean(qs); median(qs)
```

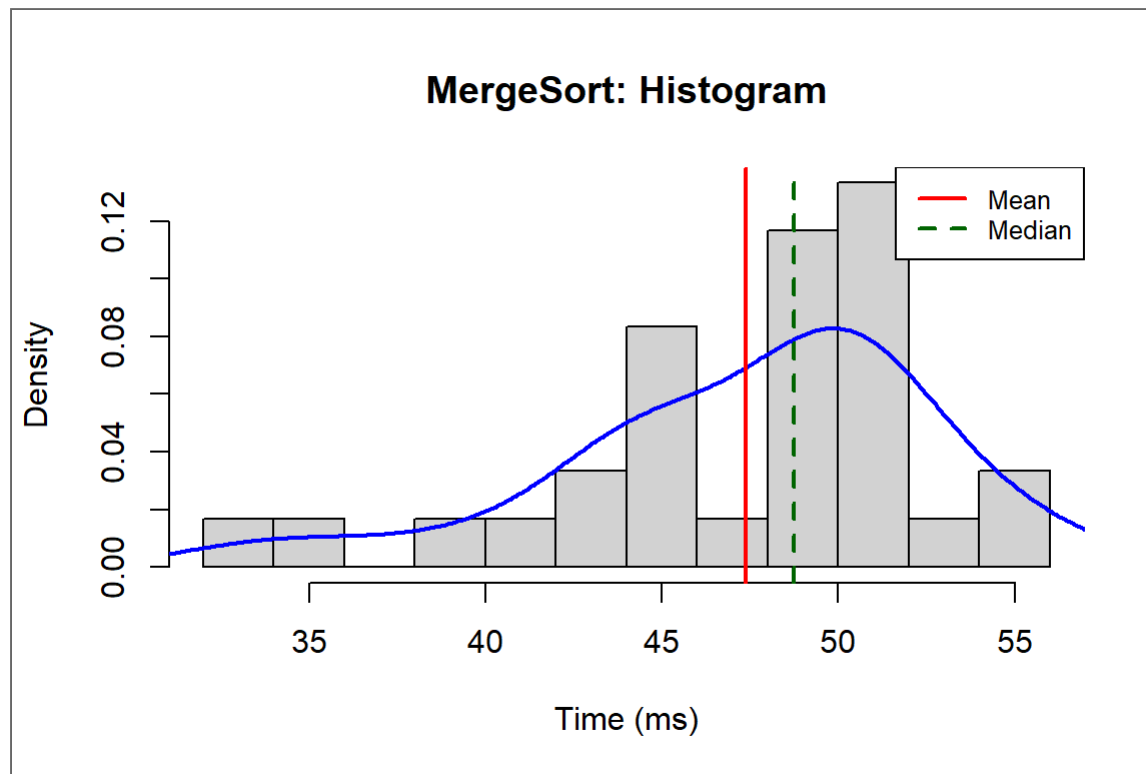
```
## [1] 52.41152
```

```
## [1] 51.39765
```

Do the mean and median differ? That can tell us about symmetry or skewness.

Visualizing the Center of Data

Let's visualize both distributions using histograms with **mean** and **median** lines.



Understanding Variability

Two datasets can have the same mean but very different **spread**. Measures of variability include:

- **Range:** max - min
- **Deviations** from the mean
- **Variance:** average squared deviation
- **Standard Deviation (SD):** square root of variance
- **IQR:** interquartile range (Q3 - Q1)

Why Standard Deviation?

Standard deviation (SD) is the most common way to quantify **how much the data spread** around the mean.

- Low SD = tightly clustered
- High SD = widely spread


```
tapply(algo_data$Time_ms, algo_data$Algorithm, sd)
```

```
## MergeSort QuickSort  
## 5.250282 7.530168
```

Boxplots: A Visual Summary

Boxplots show:

- Median (center line)
- Q1 and Q3 (box edges)
- IQR (box width)
- Outliers (dots beyond whiskers)



What Can Boxplots Tell Us?

Boxplots help you:

- Compare medians
 - Spot skewed distributions (asymmetry)
 - Detect outliers
 - Evaluate overall spread (IQR vs total range)
- Use them early and often during data exploration!

Distributions and Statistical Assumptions

Before conducting statistical tests, we must understand **how the data are distributed**.

This section covers:

- What distributions are
- The normal (bell-shaped) distribution
- How to assess normality in R
- When and why normality matters

What Is a Distribution?

A **distribution** describes how values are spread out.

Key properties:

- **Symmetry vs skewness**
- **Unimodal vs bimodal**
- **Continuous vs discrete**

Examples:

- Uniform: all values equally likely
- Normal: bell-shaped
- Skewed: lopsided toward one side

The Normal Distribution

The normal distribution (Gaussian) is:

- Symmetric
- Bell-shaped
- Defined by **mean** and **standard deviation**

In a normal distribution:

- Mean = Median = Mode
- Most values lie near the center

The Empirical Rule


The **Empirical Rule** states:

- ~68% of data lie within ± 1 SD
- ~95% within ± 2 SD
- ~99.7% within ± 3 SD

This rule is essential for:

- Understanding **confidence intervals**
- Identifying **outliers**

Visualizing the Normal Curve



Why Does Normality Matter?

Many statistical tests assume normality — especially:

- **t-tests**
- **ANOVA**
- **Regression**

We assume:

“If the sample comes from a normal distribution, the test results are valid.”

Fortunately, many tests are **robust** to small violations.

Central Limit Theorem (CLT)

The **CLT** says:

*The distribution of the sample mean becomes approximately **normal** as the sample size increases, even if the data are not normal.*

This allows us to:

- Use t-tests and CIs on non-normal data
- Rely on sampling distributions instead of raw data


CLT usually kicks in around $n \geq 30$

Graphical Checks for Normality

Three useful visual tools:

- **Histogram** — see symmetry, outliers
- **Density Plot** — smoother histogram
- **QQ Plot** — plot of quantiles vs theoretical normal values


Example: Histogram + Density Plot



QQ Plot (Quantile–Quantile Plot)

In a QQ plot:

- Points should lie on a straight line
- Curves at ends → skewness
- Strong deviations → non-normal



Formal Test: Shapiro-Wilk Test

Tests whether data are normally distributed.

```
shapiro.test(ms)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  ms  
## W = 0.93428, p-value = 0.06386
```

Interpretation:

- H_0 : Data are from a normal distribution
- H_1 : Data are not normal
- If $p > 0.05 \rightarrow$ fail to reject $H_0 \rightarrow$ data are *normal enough*

What Does “Assume Normality” Mean?

It doesn't mean your data must be perfect.

It means:

- You've **checked assumptions**
- The data are **close enough**
- You can justify using parametric tests

When in doubt:

- Use **non-parametric tests** (next section)
- Be **transparent** in your thesis

Confidence Intervals and Statistical Inference

After describing our data and checking assumptions, we're ready for inference — making conclusions about a population based on sample data. The first tool of inference is the **confidence interval**.

What Is a Confidence Interval?

A **confidence interval (CI)** is a range of values likely to contain the **true population parameter** (like a mean).

For example:

“MergeSort is faster by 4.1 ms, 95% CI [1.6, 8.4]”

This means we're 95% confident the **true difference** is between 1.6 and 8.4 ms.

How to Interpret “95% Confidence”

It **does not mean** there's a 95% chance the true value lies in one specific interval.

It means:

If we repeated this experiment 100 times, 95 of those intervals would contain the true value.

It's about the **method**, not one single outcome.

Confidence Interval Formula

The general formula for a CI is:

$$CI = \bar{x} \pm t \cdot SE$$

Where:

- \bar{x} : sample mean
- t : critical value from the t-distribution
- **SE**: standard error = s / \sqrt{n}

CI vs Standard Error

Standard Error (SE) shows how much sample means vary:

$$SE = \frac{s}{\sqrt{n}}$$

The **CI** uses SE to build a range:

- SE = “spread of sample means”
- CI = “range of plausible population values”

Computing Confidence Intervals in R

We can use `t.test()` to compute the mean and its confidence interval:

```
qs_data <- algo_data$Time_ms[algo_data$Algorithm ==  
  "QuickSort"]  
ms_data <- algo_data$Time_ms[algo_data$Algorithm ==  
  "MergeSort"]
```

```
t.test(qs_data)$conf.int
```

```
## [1] 49.59971 55.22333  
## attr(,"conf.level")  
## [1] 0.95
```

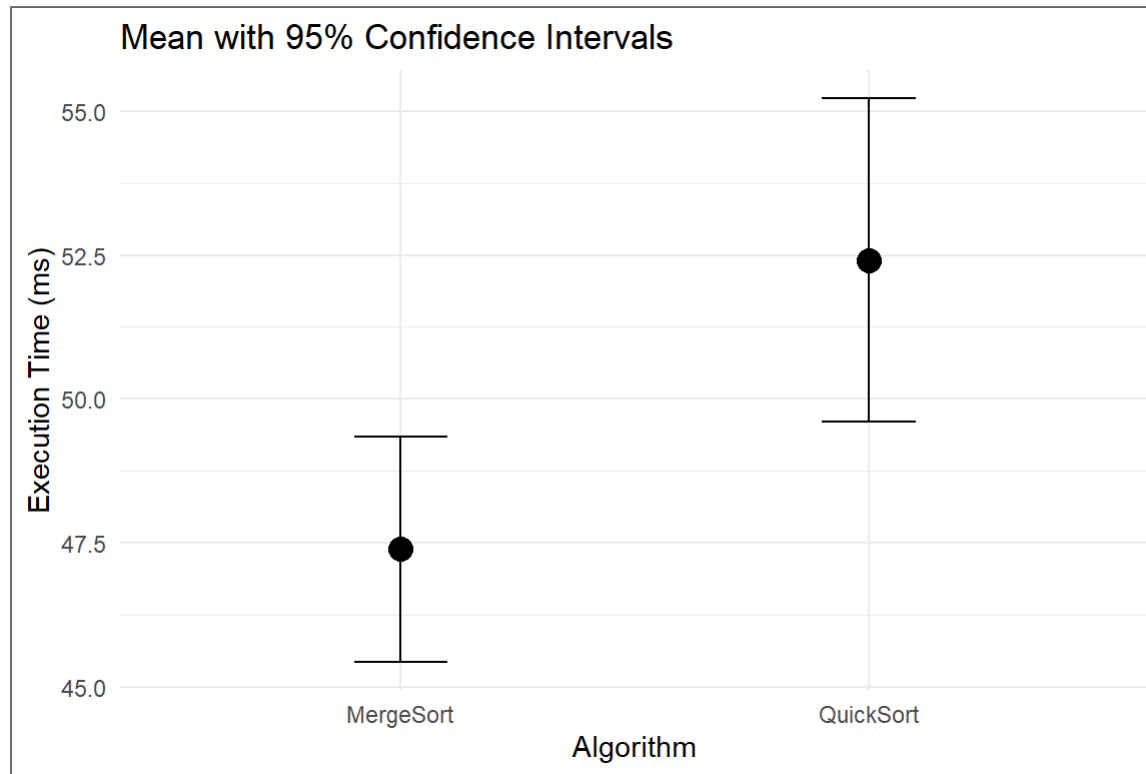
```
t.test(ms_data)$conf.int
```

```
## [1] 45.42997 49.35095  
## attr(,"conf.level")  
## [1] 0.95
```

This gives the **95% confidence interval** around the mean run-time.

Visualizing Confidence Intervals

We can plot mean \pm CI error bars for each algorithm.



CI for the Difference in Means

What if we want to compare MergeSort and QuickSort directly?

```
t.test(Time_ms ~ Algorithm, data = algo_data)$conf.int
```

```
## [1] -8.384496 -1.657631  
## attr(,"conf.level")  
## [1] 0.95
```

If the CI includes 0 → **no significant difference**

If it excludes 0 → **significant difference**

Example Interpretation

Suppose the CI is:

$[-8.38, -1.66]$

Interpretation: - We are 95% confident MergeSort is faster by **between 1.66 and 8.38 ms** - CI does not include 0 → result is **statistically significant**

Summary: Confidence Intervals

Concept	Purpose
CI	Range of plausible population values
95% CI	95 out of 100 such intervals contain the true value
CI includes 0	No significant difference
CI excludes 0	Significant difference
Narrow CI	More precision (less variability)

Confidence intervals are a powerful tool for **quantifying uncertainty** – and they naturally lead us into **hypothesis testing**.

Foundations of Hypothesis Testing

Once we have summary statistics and confidence intervals, we're ready for formal decision-making.

Hypothesis testing helps us answer questions like: > “Is MergeSort significantly faster than QuickSort?”

It is the standard method for drawing scientific conclusions from data.

What Is a Hypothesis?

A **hypothesis** is a testable statement about a population or process.

It must be:

- **Testable** — can be confirmed or refuted with data
- **Falsifiable** — you can imagine evidence against it
- **Specific** — clear about what is being measured

Types of Hypotheses

We always define two hypotheses:

- **Null Hypothesis (H_0)**

The default assumption — no difference, no effect

E.g., “The algorithms have equal mean run-time”

- **Alternative Hypothesis (H_1)**

What we want to show — there is a difference or effect

E.g., “MergeSort is faster than QuickSort”

We test whether we can reject H_0 in favor of H_1 .

What Makes a Good Hypothesis?

A strong hypothesis should: - Be based on theory or prior research

- Identify specific variables
- Be framed for statistical testing
- Use clear operational definitions

Examples:

Research Question	H ₀	H ₁
Is MergeSort faster?	$\mu_1 = \mu_2$	$\mu_1 < \mu_2$
Do students prefer UI A or B?	Equal preferences	At least one differs
Is crash rate linked to platform?	Independent	Dependent

The Logic of Hypothesis Testing

1. Assume H_0 is true
2. Collect data
3. Calculate how likely the data are if H_0 were true
4. If the data are **very unlikely**, reject H_0

This logic uses the **p-value** to quantify that likelihood.

What Is a p-value?

The **p-value** is the probability of observing your results (or more extreme), **assuming H_0 is true**.

- Small p-value → data are unlikely under H_0
- Large p-value → data are plausible under H_0

*It is **not** the probability that H_0 is true.*

Interpreting the p-value

p-value	Interpretation
≤ 0.05	Statistically significant — reject H_0
> 0.05	Not significant — fail to reject H_0

Significance level (α) is usually set at **0.05**

Type I and Type II Errors

Type	Description	Risk Symbol
Type I Error	Rejecting H_0 when it is actually true	α
Type II Error	Failing to reject H_0 when it is false	β

Reducing one type often increases the other.

We aim to **balance** them.

One-Tailed vs Two-Tailed Tests

- **Two-tailed:** Checks for **any** difference

$$H_1: \mu_1 \neq \mu_2$$

- **One-tailed:** Checks for a **specific** direction

$$H_1: \mu_1 < \mu_2$$

Use **one-tailed only** when justified by theory — not after looking at data!

The Significance Level (α)

We set a threshold to control Type I error:

- Typically: $\alpha = 0.05$
- Stricter for high-risk domains: $\alpha = 0.01$

We compare:

- If $p \leq \alpha \rightarrow$ reject H_0 (significant)
- If $p > \alpha \rightarrow$ fail to reject H_0 (not significant)

Summary: Hypothesis Testing Flow

Step	Description
1. State H_0 and H_1	Define competing hypotheses
2. Choose α	Set risk tolerance (usually 0.05)
3. Compute p-value	Based on sample and test
4. Compare p to α	Decide whether to reject H_0
5. Interpret	Report significance and effect size

This process will now guide us into **specific tests** — starting with the **t-test** for comparing means.

Comparing Means: t-Tests

When your outcome is **numeric** and you want to compare **means**, the most common tool is the **t-test**.

There are three main types:

- One-sample t-test
- Independent two-sample t-test
- Paired-sample t-test

We'll explore each with assumptions, examples, and interpretation.

One-Sample t-Test

Use when: You compare a sample mean to a known value.

Is MergeSort's mean run-time different from 50ms?

```
ms_data <- algo_data$Time_ms[algo_data$Algorithm ==  
  "MergeSort"]  
t.test(ms_data, mu = 50)
```

```
##  
## One Sample t-test  
##  
## data: ms_data  
## t = -2.7223, df = 29, p-value = 0.01085  
## alternative hypothesis: true mean is not equal to 50  
## 95 percent confidence interval:  
## 45.42997 49.35095  
## sample estimates:
```

```
## mean of x  
## 47.39046
```

- $H_0: \mu = 50$
- $H_1: \mu \neq 50$

Reject H_0 if the p-value is less than 0.05.

Independent Two-Sample t-Test

Use when: You compare two independent groups.

Is there a significant difference in mean run-time between MergeSort and QuickSort?

```
t.test(Time_ms ~ Algorithm, data = algo_data, var.equal  
      = FALSE)
```

```
##  
##  Welch Two Sample t-test  
##  
## data:  Time_ms by Algorithm  
## t = -2.9959, df = 51.806, p-value = 0.004192  
## alternative hypothesis: true difference in means  
## between group MergeSort and group QuickSort is not equal  
## to 0  
## 95 percent confidence interval:
```

```
## -8.384496 -1.657631
## sample estimates:
## mean in group MergeSort mean in group QuickSort
##                47.39046                52.41152
```

- $H_0: \mu_1 = \mu_2$
- $H_1: \mu_1 \neq \mu_2$

Welch's t-test (`var.equal = FALSE`) handles unequal variances.

Paired-Sample t-Test

Use when: You compare two values from the **same subject** (e.g., two algorithms on the same input).

Simulated Paired Data

```
set.seed(123)
paired_data <- data.frame(
  Input = 1:30,
  QuickSort = rnorm(30, 52, 6),
  MergeSort = rnorm(30, 48, 5)
)

t.test(paired_data$QuickSort, paired_data$MergeSort,
       paired = TRUE)
```

```
##  
## Paired t-test  
##  
## data: paired_data$QuickSort and paired_data$MergeSort  
## t = 2.001, df = 29, p-value = 0.05483  
## alternative hypothesis: true mean difference is not  
## equal to 0  
## 95 percent confidence interval:  
## -0.06243577 5.71380733  
## sample estimates:  
## mean difference  
## 2.825686
```

- H_0 : Median difference = 0
- H_1 : Median difference $\neq 0$

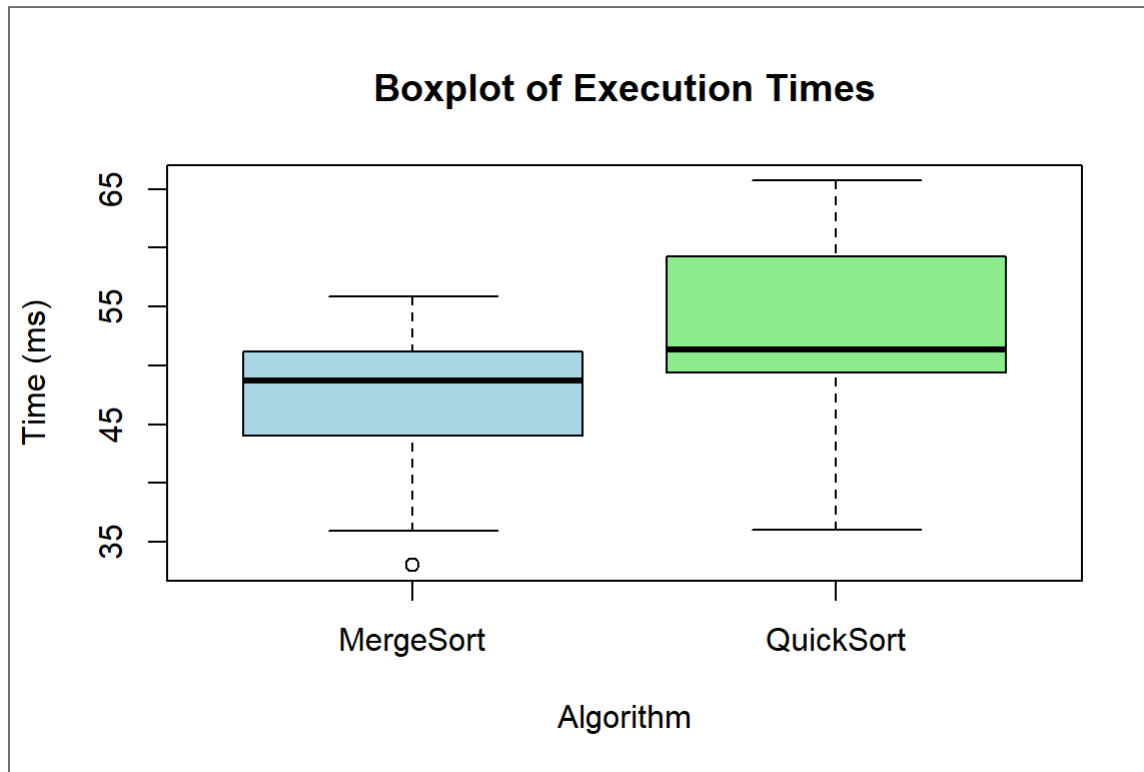
Assumptions of the t-Test

Assumption	Description
Normality	Each group's data ~ normal distribution
Equal variances	Use Welch's test if not true
Independence	Observations are not repeated
Interval-level data	Values must be numeric and continuous

Violations of these assumptions lead to **non-parametric tests** (we'll cover these later).

Visualizing Group Differences

```
boxplot(Time_ms ~ Algorithm, data = algo_data,  
        col = c("lightblue", "lightgreen"),  
        main = "Boxplot of Execution Times",  
        ylab = "Time (ms)")
```



Use boxplots and CI bars to complement test results.

Summary: Which t-Test?

Test	Use Case	Function
One-sample t-test	Sample vs fixed value	<code>t.test(x, mu = val)</code>
Two-sample t-test	Two independent groups	<code>t.test(y ~ group)</code>
Paired-sample t-test	Two related measurements	<code>t.test(x, y, paired = TRUE)</code>

Comparing 3+ Groups: ANOVA

When comparing **more than two groups**, we use **ANOVA** instead of multiple t-tests to avoid inflating the Type I error rate.

ANOVA tells us:

*“Is there **any** significant difference between group means?”*

We'll explore:

- One-way ANOVA
- Two-way ANOVA
- Post-hoc tests (Tukey HSD)
- Assumptions and visual inspection

One-Way ANOVA: Setup

Let's simulate a third algorithm: HeapSort

```
set.seed(2025)
qs <- rnorm(30, mean = 52, sd = 6)
ms <- rnorm(30, mean = 48, sd = 5)
hs <- rnorm(30, mean = 50, sd = 4.5)

algo3_data <- data.frame(
  Algorithm = rep(c("QuickSort", "MergeSort",
                    "HeapSort"), each = 30),
  Time_ms = c(qs, ms, hs)
)
```

Running One-Way ANOVA

```
anova_model <- aov(Time_ms ~ Algorithm, data =
                    algo3_data)
summary(anova_model)
```

```
##              Df Sum Sq Mean Sq F value  Pr(>F)
## Algorithm      2  408.2   204.12    7.499 0.00099 ***
## Residuals     87 2368.2    27.22
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1
```

- $H_0: \mu_1 = \mu_2 = \mu_3$ (all means are equal)
 - H_1 : At least one group differs
- If $p < 0.05 \rightarrow$ Reject $H_0 \rightarrow$ At least one mean is significantly different

Post-Hoc Testing: Tukey HSD

ANOVA tells us there is a difference, but not which groups differ.
Use **Tukey's Honest Significant Difference (HSD)** to compare all pairs:

```
TukeyHSD(anova_model)
```

```
##    Tukey multiple comparisons of means
##      95% family-wise confidence level
##
## Fit: aov(formula = Time_ms ~ Algorithm, data =
## algo3_data)
##
## $Algorithm
##
##              diff              lwr              upr
p adj
## MergeSort-HeapSort -0.9764591 -4.1886528 2.235735
##                   0.7494717
## QuickSort-HeapSort  3.9499434  0.7377497 7.162137
```



0.0118579

Look for:

- Which comparisons are significant
- Confidence intervals that do **not** include 0

Visualizing Group Differences

```
boxplot(Time_ms ~ Algorithm, data = algo3_data,  
        col = c("skyblue", "lightgreen", "lightpink"),  
        ylab = "Execution Time (ms)",  
        main = "Algorithm Comparison (3 Groups)")
```



Assumptions of ANOVA

Assumption	How to Check
Independence	Experimental design
Normality	Use QQ plot or Shapiro test per group
Equal variances	Use Bartlett or Levene test

Example: Test for Equal Variances

```
bartlett.test(Time_ms ~ Algorithm, data = algo3_data)
```

```
##  
##  Bartlett test of homogeneity of variances  
##
```

```
## data: Time_ms by Algorithm  
## Bartlett's K-squared = 1.3655, df = 2, p-value = 0.5052
```

Two-Way ANOVA: Including More Factors

Let's add **Input Size** as a second factor: Small vs Large datasets.

```
set.seed(123)
algo <- rep(c("QuickSort", "MergeSort", "HeapSort"),
            times = 30)
size <- rep(rep(c("Small", "Large"), each = 15), times
            = 3)

runtime <- rnorm(90,
                 mean = ifelse(algo == "MergeSort",
                               48,
                               ifelse(algo == "QuickSort",
                                       52, 50)) +
                 ifelse(size == "Large", 3, 0),
                 sd = 4.5)
```

Two-Way ANOVA Model

```
two_way_model <- aov(Time_ms ~ Algorithm * InputSize,
                     data = two_way_data)
summary(two_way_model)
```

```
##              Df Sum Sq Mean Sq F value
Pr(>F)
## Algorithm          2   477.3   238.67   14.485 3.94e-
06 ***
## InputSize          1   142.4   142.44    8.645
0.00424 **
## Algorithm:InputSize  2     4.9     2.45    0.149
0.86182
## Residuals         84 1384.1    16.48
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.'
0.1 ' ' 1
```

This tests:

- Main effect of **Algorithm**
- Main effect of **Input Size**
- **Interaction** between them

A significant interaction means:

The effect of Algorithm depends on Input Size

Summary: ANOVA

Test	Use Case	Tool
One-way ANOVA	3+ group means	<code>aov(y ~ group)</code>
Tukey HSD	Post-hoc pairwise	<code>TukeyHSD()</code>
Bartlett Test	Equal variances	<code>bartlett.test()</code>
Two-way ANOVA	2 factors	<code>aov(y ~ A * B)</code>

Chi-Squared Tests for Categorical Data

When your data are **categorical** (e.g., success/fail, error type, UI preference), you need a test designed for **counts**, not means.

The **Chi-squared (χ^2) test** helps answer questions like:

- Do observed results match expectations?
- Are two categorical variables related?

Two Main Types of Chi-Squared Tests

Test	Purpose
Goodness-of-Fit	Do observed counts match expected proportions?
Test of Independence	Are two categorical variables related?

We'll explore both with examples.

Chi-Squared Goodness-of-Fit

Use when: You compare observed outcomes to an expected distribution.

Example: HeapSort ran 60 times

You expect 95% Success, 5% Fail.

You observed: 54 Success, 6 Fail

```
observed <- c(Success = 54, Fail = 6)
expected_probs <- c(0.95, 0.05)
chisq.test(x = observed, p = expected_probs)
```

```
## Warning in chisq.test(x = observed, p =
expected_probs): Chi-squared
## approximation may be incorrect
```

```
##
## Chi-squared test for given probabilities
```

```
##
```

```
## data:  observed
```

```
## X-squared = 3.1579, df = 1, p-value = 0.07556
```

Interpreting Goodness-of-Fit

- H_0 : The observed distribution matches the expected
- H_1 : The distributions differ
- If $p < 0.05 \rightarrow$ significant \rightarrow reject H_0

This helps check whether results conform to a known or ideal pattern.

Chi-Squared Test of Independence

Use when: You want to know if two categorical variables are associated.

Example: Is success rate dependent on the algorithm used?

Simulated Example (Success/Failure by Algorithm)

```
## Warning in chisq.test(tab): Chi-squared approximation  
may be incorrect
```

```
##  
## Pearson's Chi-squared test  
##  
## data: tab  
## X-squared = 9.6041, df = 2, p-value = 0.008213
```

Interpreting the Test of Independence

- H_0 : The variables are independent (algorithm does not affect outcome)
 - H_1 : The variables are dependent (algorithm affects outcome)
- If $p < 0.05$, we reject H_0 and conclude that the outcome **depends** on the algorithm.

Assumptions of the Chi-Squared Test

Assumption	How to Handle
Independent observations	No repeated measurements
Expected count ≥ 5 in each cell	Use Fisher's Exact Test if violated
Sufficient total sample size	Larger n = better approximation

When Expected Counts Are Too Small

If any expected cell count is **less than 5**, consider:

- Combining similar categories
- Using **Fisher's Exact Test** (especially for 2x2 tables)

```
fisher.test(tab)  # if dimensions are small
```

```
##  
## Fisher's Exact Test for Count Data  
##  
## data:  tab  
## p-value = 0.01512  
## alternative hypothesis: two.sided
```

Summary: Chi-Squared Tests

Test Type	Use Case	R Function
Goodness-of-fit	One variable vs expected proportions	<code>chisq.test(x, p)</code>
Test of independence	Two categorical variables	<code>chisq.test(table)</code>
Small samples	Any 2x2 table	<code>fisher.test()</code>

These tests are essential when your research deals with:

- User preferences
- Pass/fail outcomes
- Classification results
- Any system behavior encoded as **categories**

Correlation Analysis

When working with **two numeric variables**, we often ask:

“As one variable increases, does the other also increase (or decrease)?”

Correlation analysis helps answer this by measuring:

- The **strength** of the relationship
- The **direction** (positive or negative)

Types of Correlation

Method	Use When
Pearson	Linear relationship, numeric data, normality
Spearman	Ordinal data or non-linear monotonic relationship
Kendall (optional)	Small samples, tied ranks

We'll focus on Pearson and Spearman.

Pearson Correlation Coefficient (r)

Used for:

- Two continuous, normally distributed variables
- Linear relationships

Values range from:

- +1: perfect positive correlation
- 0: no correlation
- -1: perfect negative correlation

```
set.seed(1)
input_size <- seq(100, 1000, length.out = 50)
qs_runtime <- 0.05 * input_size + rnorm(50, mean = 0, sd
    = 10)

cor(input_size, qs_runtime) # Pearson correlation
```

```
## [1] 0.8478924
```

Spearman Rank Correlation

Used for:

- Ordinal or **non-linear monotonic** relationships
- Data with **outliers** or **non-normality**


```
cor(input_size, qs_runtime, method = "spearman")
```

```
## [1] 0.8690036
```

Compares **ranks** instead of raw values.

Scatterplot: Input Size vs Run Time

Always visualize your correlations — outliers and non-linear trends are not visible from `cor()` alone.



Correlation Matrix (Multiple Variables)

You can quickly assess relationships between several numeric variables using a correlation matrix:

```
df <- data.frame(  
  Size = input_size,  
  QuickSort = qs_runtime,  
  Noise = rnorm(50),  
  Quadratic = input_size^2 + rnorm(50, sd = 100)  
)  
  
round(cor(df), 2)
```

```
##           Size QuickSort Noise Quadratic  
## Size      1.00      0.85 -0.08      0.98  
## QuickSort 0.85      1.00 -0.09      0.84
```

## Noise	-0.08	-0.09	1.00	-0.07
## Quadratic	0.98	0.84	-0.07	1.00

Use this to explore pairwise relationships at a glance.

Correlation \neq Causation

Just because two variables move together does **not** mean one causes the other.

Correlation may result from:

- *Coincidence*
- *Confounding (a third variable affects both)*
- *Reverse causality*

Always interpret correlations with **context and caution**.

Summary: Correlation Analysis

Concept	Pearson	Spearman
Measures	Linear association	Monotonic association
Sensitive to	Outliers, skew	Less sensitive
Use with	Numeric, normal data	Ranked or skewed data
Output	-1 to +1 (strength & direction)	Same scale, different method

Correlation is a **powerful exploratory tool**, but should be followed by **modeling or experimental design** when aiming for causal claims.

Non-Parametric Alternatives

Sometimes, your data:

- Isn't **normally distributed**
- Has **small sample sizes**
- Includes **outliers** or **ranks** instead of true numeric values

In these cases, we turn to **non-parametric tests**.

*These tests don't rely on distributional assumptions and operate on **ranks** instead of raw data.*

When to Use Non-Parametric Tests

Use them when:

- Normality is violated
- You're working with **ordinal** (ranked) data
- Your sample size is too small for the Central Limit Theorem to apply
- Outliers distort the mean

Wilcoxon Signed-Rank Test

Alternative to: Paired t-test

Use when: - You have two related groups
- The differences are not normally distributed

```
qs <- rnorm(30, mean = 52, sd = 7)
ms <- rnorm(30, mean = 48, sd = 7)

wilcox.test(qs, ms, paired = TRUE)
```

```
##
##  Wilcoxon signed rank exact test
##
## data:  qs and ms
## V = 358, p-value = 0.008705
## alternative hypothesis: true location shift is not
## equal to 0
```


- H_0 : Median difference = 0
- H_1 : Median difference \neq 0

Mann–Whitney U Test (Wilcoxon Rank-Sum)

Alternative to: Independent two-sample t-test

Use when:

- Comparing two independent groups
- The data are skewed or ordinal

```
group <- rep(c("QuickSort", "MergeSort"), each = 30)
times <- c(qs, ms)

wilcox.test(times ~ group)
```

```
##
##  Wilcoxon rank sum exact test
##
## data:  times by group
## W = 283, p-value = 0.01307
```

```
## alternative hypothesis: true location shift is not  
equal to 0
```

- Compares whether one group tends to rank higher than the other

Kruskal–Wallis Test

Alternative to: **One-way ANOVA**

Use when: - You want to compare **3 or more groups**
- Data is not normally distributed

```
group <- rep(c("QuickSort", "MergeSort", "HeapSort"),  
            each = 30)  
times <- c(  
  rnorm(30, mean = 52, sd = 6),  
  rnorm(30, mean = 48, sd = 6),  
  rnorm(30, mean = 50, sd = 6)  
)  
  
kruskal.test(times ~ group)
```

```
##  
##  Kruskal-Wallis rank sum test  
##
```

```
## data:  times by group  
## Kruskal-Wallis chi-squared = 3.4113, df = 2, p-value  
= 0.1817
```

If significant → follow up with pairwise Wilcoxon tests.

Summary: Parametric vs Non-Parametric

Goal	Parametric Test	Non-Parametric Alternative
Compare one sample to a value	One-sample t-test	(rarely needed)
Compare two independent groups	Two-sample t-test	Mann-Whitney U
Compare two paired groups	Paired t-test	Wilcoxon signed-rank
Compare 3+ groups	One-way ANOVA	Kruskal-Wallis

When to Prefer Non-Parametric Tests

Choose non-parametric tests if:

- The **normality assumption** fails
- You're analyzing **ordinal/ranked data**
- You want a method **robust to outliers**

They're less powerful than parametric tests **when assumptions are met**, but much more **reliable** when assumptions are **violated**.

Choosing the Right Statistical Test

Goal	Data	Groups	Assumptions Met?	Test
Compare to known value	Numeric	1	✓	One-sample t-test
Compare two groups	Numeric	2 (independent)	✓	Two-sample t-test
Compare two paired sets	Numeric	2 (paired)	✓	Paired t-test
Compare 3+ groups	Numeric	3+	✓	ANOVA
Compare ranks/medians	Ordinal / skewed	2	✗	Mann-Whitney U / Wilcoxon

Goal	Data	Groups	Assumptions Met?	Test
Compare 3+ medians	Ordinal / skewed	3+	✗	Kruskal-Wallis
Compare counts to expected	Categorical	1 variable	–	Chi-squared (Goodness-of-Fit)
Association between categories	Categorical	2 variables	–	Chi-squared (Independence)
Correlation (linear)	Numeric	2 vars	✓	Pearson
Correlation (ranked)	Ordinal / skewed	2 vars	✗	Spearman